

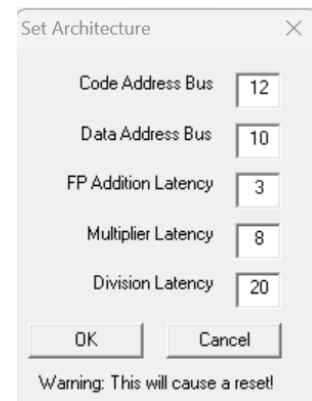
Laboratory
2

Expected delivery of lab_02.zip must include:

- **program_1.s** and **program_2.s**
- This file, filled with information and possibly compiled in a pdf format.

Please, configure the winMIPS64 simulator with the *Base Configuration* provided in the following (*in italics not user controllable configuration*):

- Code address bus: 12
- Data address bus: 10
- Pipelined FP arithmetic unit (latency): 3 stages
- Pipelined multiplier unit (latency): 8 stages
- divider unit (latency): not pipelined unit, 20 clock cycles
- Forwarding is enabled
- Branch prediction is disabled
- Branch delay slot is disabled
- *Integer ALU: 1 clock cycle*
- *Data memory: 1 clock cycle*
- *Branch delay slot: 1 clock cycle.*



- 1) Write an assembly program (**program_1.s**) for the *winMIPS64* architecture described before able to implement the following piece of code described at high-level:

```
for (i = 0; i < 64; i++){
    v5[i] = ((v1[i]* v2[i]) + v3[i])+v4[i];
    v6[i] = v5[i]/(v4[i]+v1[i]);
    v7[i] = v6[i]*(v2[i]+v3[i]);
}
```

Assume that the vectors `v1[]`, `v2[]`, `v3[]`, and `v4[]` are allocated previously in memory and contain 64 double precision **floating point** values; assume also that `v1[]` and `v4[]` do not contain 0 values. Additionally, the vectors `v5[]`, `v6[]`, `v7[]` are empty vectors also allocated in memory.

Calculate the data memory footprint of your program:

Data	Number of Bytes
V1	512
V2	512
V3	512
V4	512
V5	512
V6	512
V7	512
Total	3584

Are there any issues? Yes, where and why? No ? Do you need to change something?

Your answer: Non c'è abbastanza memoria; si aumenta il data bus a 12 bit, per ottenere uno spazio di indirizzamento di 4K bytes, adatti a contenere i 3584 bytes necessari.

ATTENTION: winMIPS64 has a limitation due to the underlying software. There is a limitation in the string length when declaring a vector. Split the vectors elements in multiple lines (it also increases the readability) .

Example: my_fancy_vector: .byte 4, 5 ,7, 8
.byte 5,77, 8
.byte

a. Compute the CPU performance equation (CPU time) of the previous program following the next directions, assume a clock frequency of 1MHz:

$$\text{CPU time} = \left(\sum_{i=1}^n \text{CPI}_i \times \text{IC}_i \right) \times \text{Clock cycle period}$$

- Count manually, the number of the different instructions (IC_i) executed in the program
- Assume that the CPI_i for every type of instructions equals the number of clock cycles in the instruction EXE stage, for example:
 - integer instructions $\text{CPI} = 1$
 - LD/SD instructions $\text{CPI} = 1$
 - FP MUL instructions $\text{CPI} = 8$
 - FP DIV instructions $\text{CPI} = 20$
 - ...

b. Compute by hand again the CPU performance equation assuming that you can improve the FP Multiplier or the FP Divider by speeding up by 2 only one of the units at a time:

- Pipelined FP multiplier unit (latency): $8 \rightarrow 4$ stages
Or
- FP Divider unit (latency): not pipelined unit, $20 \rightarrow 10$ clock cycles

Table 1: CPU time by hand

	CPU Time initial (a)	CPU Time (b – MUL speed up)	CPU Time (b – DIV speed up)
program_1.s	3.782 ms	3.270 ms	3.142 ms

c. Using the simulator calculate again the CPU time and complete the following table:

Table 2: CPU time using the simulator

	CPU Time initial (a)	CPU Time (b – MUL speed up)	CPU Time (b – DIV speed up)
program_1.S	3.468 ms	2.956 ms	2.828 ms

Are there any difference? If yes, where and why? If Not, provide some comments in the following:

Your answer: Sì, sono presenti differenze.
 Il simulatore opera con pipelining e forwarding: nonostante la presenza di stalli (dovuti ad esempio alle dipendenze di dato o agli accessi concorrenti allo stesso modulo), la possibilità di parallelizzare le operazioni consente di terminare l'esecuzione del programma rispetto ai calcoli effettuati manualmente. In questo caso, non si considerava la possibilità di stalli, ma nemmeno la parallelizzazione delle operazioni.

d. Using the simulator and the *Base Configuration*, disable the Forwarding option and compute how many clock cycles the program takes to execute.

Table 3: forwarding disabled

	Number of clock cycles	IPC (Instructions Per Clock)
program_1.S	4749	0.244

Enable one at a time the **optimization features** that were initially disabled and collect statistics to fill the following table (fill all required data in the table before exporting this file to pdf format to be delivered).

Table 4: **Program performance for different processor configurations**

Program	Forwarding		Branch Target Buffer		Delay Slot		Forwarding + Branch Target Buffer	
	IPC	CC	IPC	CC	IPC	CC	IPC	CC
Program_1.S	0.334	3468	0.247	4688	0.244	4749	0.340	3407

- 2) Using the WinMIPS64 simulator, validate experimentally the Amdahl's law, defined as follows:

$$\text{speedup}_{\text{overall}} = \frac{\text{execution time}_{\text{old}}}{\text{execution time}_{\text{new}}} = \frac{1}{(1 - \text{fraction}_{\text{enhanced}}) + \frac{\text{fraction}_{\text{enhanced}}}{\text{speedup}_{\text{enhanced}}}}$$

- a. Using the program developed before: **program_1.s**
- b. Modify the processor architectural parameters related with multicycle instructions (Menu→Configure→Architecture) in the following way:
 - 1) Configuration 1
 - Starting from the *Base Configuration*, change only the FP addition latency to 6
 - 2) Configuration 2
 - Starting from the *Base Configuration*, change only the Multiplier latency to 4
 - 3) Configuration 3
 - Starting from the *Base Configuration*, change only the division latency to 10

Compute by hand (using the Amdahl's Law) and using the simulator the speed-up for any one of the previous processor configurations. Compare the obtained results and complete the following table.

Table 5: **program_1.s** speed-up computed by hand and by simulation

Proc. Config.	Base config. [c.c.]	Config. 1	Config. 2	Config. 3
Speed-up comp.				
By hand	3782	0.80	1.37	1.20
By simulation	3468	0.90	1.17	1.23

- 3) Write an assembly program (**program_2.s**) for the winMIPS64 architecture able to compute the output (y) of a **neural computation** (see the Fig. below):


$$x = \sum_{j=0}^{K-1} i_j * w_j + b$$

$$y = f(x)$$


where, to prevent the propagation of NaN (Not a Number), the activation function f is defined as:

$$f(x) = \begin{cases} 0, & \text{if the exponent part of } x \text{ is equal to } 0x7ff \\ x, & \text{otherwise} \end{cases}$$

Assume the vectors i and w respectively store the inputs entering the neuron and the weights of the connections. They contain $K=30$ double precision **floating point** elements. Assume that b is a double precision **floating point** constant and is equal to $0xab$, and y is a double precision **floating point** value stored in memory. Compute y .



Below is reported the encoding of IEEE 754 double-precision binary floating-point format:



Given the *Base Configuration*, run your program and extract the following information.

	Number of clock cycles	Total Instructions	CPI (Clock per Instructions)
program_2.S	358	257	1.393